

AgentScore

Business Case for Tricentis Leadership

"Point it at a trace, get back a defensible answer."

● September 2026

The Market Signal

"Agents are being built faster than anyone can validate them."

5% → 40%

Task-specific agents as share of enterprise apps, 2025 → end of 2026 (Gartner)

150,000+

Agents projected across the Fortune 500 within two years — avg. enterprise already runs ~12, +67% projected

\$7.8B → \$11.5B

AI agent market size, 2025 → 2026 (49.6% CAGR through 2033)

1 in 5

Companies with a mature governance model for autonomous agents — 80% deploy without one

40%+

Of agentic AI projects Gartner projects will be canceled by end of 2027 due to cost, unclear value, and inadequate risk controls - the exact failure mode AgentScore exists to prevent

Demand Isn't Hypothetical

Inbound interest, last six weeks

74

Inbound-interested companies



9

Scheduled demos



1


Pending beta — Accenture, on behalf of Meta



Sourced from Gartner and Digital Applied.

The Category Is Consolidating Fast

— and validating our model along the way.



Galileo → acquired by Cisco, folded into Splunk Observability (May 2026). Langfuse → acquired by ClickHouse (Jan 2026).

Combined: ~\$258.5M raised across these four — two already acquired (prices undisclosed).

WHAT THIS VALIDATES

Every eval-first competitor (Braintrust, PromptLayer, Humanloop, Patronus) already meters “scoring/eval runs” as a distinct, premium-priced line, separate from ingestion.

Validates our own billing direction.

No competitor leads with zero-setup automatic grading straight from an OTEL trace.

Our clearest structural differentiator today.

Why Tricentis Should Solve This

Being right about the market isn't the same as being the right company to act on it — but here we are.

01

TRUST HEAD START

Quality Engineering, From Code to Agents

Users already trust Tricentis for Tosca and qTest — AgentScore isn't a new tool to evaluate, it's the same platform extended from code to the agents that code powers.

02

DISTRIBUTION ADVANTAGE

A Cost Edge No Outside Vendor Can Match

AI Workspace agents auto-provision a tenant and get graded with zero setup — every AIW customer is a warm lead, not a cold integration decision.

03

THE CLOCK IS RUNNING

The Market Is Deciding Right Now

Braintrust, Galileo, and Langfuse all moved — raised, acquired, or folded in — within months of each other. Every quarter we wait, someone else becomes the default answer to “is this agent safe to ship?”

What AgentScore Is

"Point it at a trace, get back a defensible answer."

1

Any Trace, Any Agent

Ingests an agent's OTel trace — any agent, any vendor, any orchestration framework, no proprietary SDK required.

2

A Verdict, Not a Pass/Fail

Returns a 0-100 score and a ship/warn/block verdict.

3

Root-Cause Attribution

Down to the span, paired with a suggested fix.

4

60+ Evals, 11 Dimensions

Auto-selected via zero-setup profiling — the buyer doesn't need to know what to measure going in.

5

Tiered Evidence Model

Deterministic checks, judge + confidence intervals, golden-dataset back-testing — audit-ready proof, not a black box.

6

Automatic Inside AI Workspace

AI Workspace agents get all of this free, without any extra setup, the moment the agent exists.

TARGET PERSONA

Who Actually Buys This

Ranked by priority — who reacts to the score, and why.

Persona	Why They Care	Priority
QA/quality leader at an agent-deploying enterprise	Owns “is this agent safe to ship” but has no tooling built for agent behavior — wants a defensible score and root-cause attribution.	Primary
AI/platform program owner at scale	Owns dozens of agents shipping to prod with fragmented ownership — reacts most strongly to the single ship/warn/block verdict.	Primary
SI/consulting partner building agents for clients	Needs third-party evidence of quality to close their own client deals — a resale motion, not just internal use.	Secondary
Regulated-industry buyer	Compliance/audit sensitivity makes “how do you know it's accurate” a gating question before production.	Secondary

TARGET ICP

What Tells You an Account Is Worth Targeting

If you're a sales rep, here's what to look for in an account.

- Already running agents in AI Workspace → zero-setup, warm lead, not a cold integration decision.
- Running multiple agents or workflows, not just one → more surface area, more manual grading pain.
- Risk-driven signals: talk of governance/spend concerns, “how do we know this is safe” before “how do we ship more,” or a token bill that's stopped tracking cleanly to work done.
- Efficiency-driven signals: already grading agents themselves — ad hoc scripts, spreadsheet rubrics, a homegrown eval harness — and want to cut the manual cost, or want more rigor than what they've built.
- Running external agents/tools/MCPs alongside AI Workspace (GitHub Copilot, Copilot Studio, homegrown frameworks, other vendors, MCP servers) that they can't currently score side-by-side with their sanctioned agents.

Bonus signal: no dedicated AI engineers or existing eval tooling in-house → AgentScore isn't competing against a build option, it's the only realistic path to a rigorous eval.

Wolters Kluwer is the cleanest real match today (risk-driven, AIW + external agents). Tritusa fits 3 of 4. A hypothesis, not yet a validated segment.

VALUE PROP

Why Customers Will Pay

Five reasons buyers say yes — and the economics behind them.

- A defensible verdict instead of a gut call — replaces “seems fine to me” with a number leadership can act on.
- Root-cause attribution, not just a score — turns debugging into a fix, not a re-run.
- A shared yardstick across fragmented ownership — the fleet-level ask we're already hearing (Wolters Kluwer, Freddie Mac).
- Coverage without in-house AI expertise — zero-setup auto-profiling means the buyer doesn't need to already know what to measure.
- Third-party evidence SIs can resell — a resale motion for partners like Tritusa, Capgemini, Accenture, Xebia.

OVERALL VALUE

2.3 / qtr

AI-driven errors/qtr, \$50K–\$2.1M each

\$14K/yr

per employee/yr verifying AI outputs

The billing model is already proven by the category — eval-first competitors meter scoring runs at a premium, separate from ingestion. A subscription only has to prevent a fraction of one incident, reclaim a sliver of wasted token spend, or claw back a few hours of eval-authoring time to pay for itself many times over.

ROI METRICS

The ROI Metrics We Can Measure

Three candidate metrics — none validated by a paying customer yet.

TIME SAVED ON EVALS

??? hrs/agent/week

Baseline unmeasured — no pilot customer has tracked eval-authoring time yet

40%

Hypothesis: time saved per agent per week if AgentScore auto-generates evals vs. manual authoring

\$???/yr

Per-employee value reclaimed if the hypothesis holds

% OF TOKEN SPEND REDUCTION

???

No real customer token bill to benchmark against yet

40–60%

Gartner: agentic token spend that contributes nothing to the answer. Hypothesis: AgentScore can flag and reduce this share.

Bill vs. flag

Compare AgentScore-flagged low-value spend against a pilot customer's actual token bill, before vs. after AgentScore, net of usage growth.

REDUCTION IN AI-RELATED INCIDENTS

???

Incident counts tied to agent behavior (bad outputs, safety violations, compliance failures)

20 %

Hypothesis: continuous scoring catches regressions before they reach production, reducing incident volume.

Before vs. after

Compare AI-related incident counts at a pilot account before vs. after AgentScore, over a matched time window.

STATUS: All three are hypotheses, not yet validated by a paying customer — each needs before/after pilot data before it goes in front of a buyer as a claim.