



AgentScore Competitive Positioning

Where we stand in the agentic eval market

● SEPTEMBER 2026

We are the **verdict layer** competitors aren't building.

Zero-setup automatic grading - no SDK, no test authoring, no labeled data required.

OTel-in, one defensible 0-100 score out.

The agentic eval market splits into four camps



Observability / tracing-first

Focus on capturing traces, then add evals on top.

Langfuse · LangSmith · Arize · W&B Weave



Eval / scoring-first

Quality gates and judge models are the main offering.

Braintrust · Galileo · Confident AI / DeepEval · Comet Opik



Simulation-first

Generate synthetic users and scenarios to stress-test before prod.

Arato · Patronus · Maxim



AI TRiSM / runtime guardrails

Intercept and gate agent outputs at runtime rather than grade behavior.

int4 TrustGate · NeuralTrust · Guardrails AI

Our closest neighbors: Galileo and Braintrust

Galileo	Braintrust
VALUE PROP Continuous eval + guardrails via proprietary Luna-2 SLM judges - sub-200ms, ~98% cheaper than frontier LLM-as-judge.	VALUE PROP Eval-first observability: capture traces, score with scorers, turn failures into CI.
PRICING Free (5k traces/mo); Pro \$100/mo; Enterprise custom.	PRICING Free (1GB/10k scores); Pro \$249/mo; Enterprise custom (SSO, RBAC, self-host).
EVAL / SCORING Continuous production scoring; LLM-as-judge + SLM judges; agent/safety metrics.	EVAL / SCORING Real-time + batch scoring; configurable LLM-as-judge; "Loop" AI assistant.
KEY SIGNAL ~\$68M raised. Acquired by Cisco (May 2026) → Splunk. Customers: Verizon, Comcast, HP.	KEY SIGNAL \$36M Series A + \$80M Series B (~\$800M val.). Customers: Notion, Stripe, Vercel, Ramp.

Bottom line: neither leads with "zero-setup automatic grading."

How AgentScore compares - scoring and eval capability

✔ shipped / mature 🟡 partial, gated, or in progress ✖ absent n/p not public

Dimension	AgentScore	Langfuse	Braintrust	Arize	Galileo	LangSmith	W&B Weave	Arato	Patronus	int4 TrustGate [†]
Core stance	Auto-grade real behavior, zero setup	OSS observability	Eval-first + CI	Observability + eval	Continuous scoring + guardrails	Tracing for LangChain	Drop-in trace + eval	Sim-first testing	Judge models + sim	SAP write-txn guardrail
Trace/run scoring	✔ (20-trace floor)	✔	✔	✔	✔	✔	✔	✔ (sim)	✔	🟡 document-level, not trace
Session-based scoring	✔	✔	✔	✔ (session-level)	✔	✔	✔	✔ (multi-turn)	✔	✖
LLM-as-judge	✔ (G-Eval, Hybrid)	✔	✔	✔	✔ + SLM judges	✔	✔	🟡 n/p	✔ (Lynx/GLIDER)	✖ deterministic
Multi-model judge catalog	✔ 7 models (Bedrock)	✔	✔	✔	✔	✔	✔	n/p	✔	✖ (n/a)
Customer-selectable judge	✔	✔	✔	✔	✔	✔	✔	n/p	✔	✖ (n/a)
Zero-config / auto-profile	✔ differentiated	✖	✖	✖	✖	✖	✖	🟡	✖	✖
Simulation / synthetic users	✖	✖	✖	✖	✖	✖	✖	✔	✔	🟡 prod replay (Prove)
Guardrails / runtime protect	✖	✖	🟡	🟡	✔	✖	✖	✖	🟡	✔ core

[†] int4 TrustGate is AI TRiSM / SAP runtime guardrails, not a general agent-quality evaluator - included for reference, not as a like-for-like rival.

How AgentScore compares - enterprise, deployment, and ecosystem


✔ shipped / mature 🟡 partial, gated, or in progress ✖ absent n/p not public

Dimension	AgentScore	Langfuse	Braintrust	Arize	Galileo	LangSmith	W&B Weave	Arato	Patronus	int4 TrustGate [†]
RBAC	✖ not built	🟡 paid	✔ Ent	✔ Ent	✔ Pro+	🟡 Ent	🟡 Ent	n/p	🟡	n/p
SSO/SAML	✖ not built	🟡 add-on	✔ Ent	✔ Ent	✔ Ent	✔ Ent	✔ Ent	n/p	🟡 OIDC	n/p
API key mgmt	🟡 ingest keys only	✔	✔	✔	✔	✔	✔	n/p	✔	n/p
SDK breadth	✖ none (OTel only); Read API reserved	Py, JS	6 langs	Py/TS/Java	Py/TS	Py/JS(+)	Py/TS	API+SDK	Py/TS/Java	✖ SAP-only, no SDK
OTel-native ingest	✔ (only path)	✔	✔	✔	✔	✔	✔	🟡	✔	✖ (SAP proxy)
Self-host / on-prem	✖ cloud-only (VPN-gated)	✔	✔ Ent	✔	✔	✔ Ent	✔ Ent	n/p	✔	n/p (in-line ERP proxy)
Deployment maturity	🟡 internal beta, 0 ext customers	GA	GA	GA	GA	GA	GA	early	GA	🟡 new (~H2 2026)
Ecosystem hook	✔ AIW, any OTel-compatible agent	LangChain-agnostic	broad	broad	Splunk/Cisco	LangChain	W&B/CoreWeave	any stack	Datadog	✔ SAP (Joule/Agentforce)

[†] int4 TrustGate is AI TRiSM / SAP runtime guardrails, not a general agent-quality evaluator - included for reference, not as a like-for-like rival.

Zero-setup automatic grading


- No SDK, no labeled data, no test authoring, no AI expertise needed
- Auto agent recognition plus automatic profile-fitting
- None of the other competitors we profiled offer this combination
- Alternatives all require instrumentation choices, dataset curation, or scorer configuration - or some mix of all three



01

One defensible answer, with evidence

- A single 0-100 score
- Five-band verdict, spanning ship, review, and don't-ship outcomes - with a confidence interval
- Every number is evidence-backed
- Built as a decision-maker artifact, not a dashboard



SHIP · ±4 CONFIDENCE

Grading traces too large for a model's context window

- Grades 30-50 MB traces - large enough to exceed a typical LLM context window
- The agentic orchestrator fetches only the relevant parts - about 0.2% of trace data in testing
- A concrete, demoable edge
- None of the profiled products advertise this capability

30-50 MB trace, exceeds a typical LLM context window

0.2% of trace data fetched to grade it

03

The TAIS and AI Workspace ecosystem hook

- AIW agents auto-provision a tenant and get graded with zero setup
- A distribution channel inside Tricentis that no external vendor can replicate
- This is the moat: no outside vendor can reach agents through TAIS / AI Workspace the way we can



04

No enterprise auth

- RBAC and SSO/SAML are basic requirements for most enterprise customers
- AgentScore has separate superadmin and customer roles today, but they are extremely limited and do not provide true RBAC
- For a Tricentis-branded enterprise product, this is table stakes - we cannot sell without it

✕ AgentScore: no true RBAC

● Langfuse: RBAC (paid)

✓ Braintrust: RBAC (Ent)

✓ Arize: RBAC (Ent)

✓ Galileo: RBAC (Pro+)

● LangSmith: RBAC (Ent)

● W&B Weave: RBAC (Ent)

01

Cloud-only

- Langfuse, Arize, Galileo, Braintrust, and Patronus all offer self-host, VPC, or on-prem deployment
- AgentScore today is cloud-only and VPN-gated
- A good portion of our users will require on-prem or self-host capabilities

✗ AgentScore: cloud-only, VPN-gated

✓ Langfuse: self-host

✓ Arize: self-host

✓ Galileo: SaaS/VPC/on-prem

✓ Braintrust: self-host (Ent)

✓ Patronus: self-host

02

No SDK or read API

- OTel-only ingest is elegant for zero-setup
- The reserved ck_Read API and Python SDK are unimplemented
- Competitors expose programmatic read and eval access
- Limits CI integration and data-out use cases

✓ OTel ingest

The only ingestion path today - elegant, zero-setup, live in production.

✗ Read API / SDK

The reserved ck_Read API and Python SDK remain unimplemented - no CI or data-out path yet.

03

No simulation capability

- The hottest adjacent subsegment - Arato, Patronus, Maxim - generates synthetic users to test pre-prod
- AgentScore only grades what already happened
- Not our lane today
- A strategic hole if buyers consolidate around "test + grade" as one purchase

✓ Grade what happened

We score real production traces after the fact - credible, but reactive.


✗ Simulate before it happens

Arato, Patronus, and Maxim generate synthetic users pre-prod - not our lane today.

04

No runtime guardrails

- Galileo's inline Protect runs on proprietary SLM judges
- Those judges score sub-200ms and roughly 98% cheaper than frontier LLM-as-judge calls
- Our frontier-Bedrock judge stack does not meet that cost/latency bar for high-volume scoring



05

Maturity

- Internal alpha moving to beta - no live external customers yet
- Agentic judge cutover incomplete: 2-3 PRs still pending
- The validation gate was redefined same-day to pass - a process concern, not just a feature gap
- Chat assistant and recommendations are built but dark, not yet live
- Every competitor in the comparison table is GA



06

Where AgentScore sits

We do not win a feature-checklist war today. We win by being **the grading layer no one else is**, delivered where no one else can reach **(inside TAIS)**.

On the core scoring loop the technology is credible and in some ways differentiated (zero-setup, large-trace, evidence-backed verdicts). But with no external customers, no enterprise auth, cloud-only VPN-gated access, and no public pricing, AgentScore is **not yet a market participant** - it is a strong internal alpha with a defensible wedge.

\$50-125M

Competitor funding rounds

3 acquired

ClickHouse · Cisco · CoreWeave

GA

Every competitor profiled, with marquee logos

0

AgentScore: external customers

Lead with the features competitors can't easily copy

- "Automatic grading with zero setup, inside your Tricentis stack." No other competitor offers this
- Don't position as another observability/eval dashboard - we lose that comparison to incumbents
- Position as the verdict layer (ship / don't ship + evidence) that plugs directly into AI Workspace with no instrumentation
- Also scores customers' existing/external agents
- Make large-trace grading and the evidence-backed 0-100 score the demo centerpiece

Observability dashboards	Langfuse · LangSmith · Arize
Eval frameworks	Braintrust · Galileo
The verdict layer	ship / don't ship + evidence - AgentScore



Knowing what to measure is the hardest part of evals

- Plenty of open-source eval frameworks already exist
- Actually evaluating an agent isn't the hard part
- Knowing what needs to be evaluated, and ensuring every relevant aspect is measured, is the hard part
- This is what gives users confidence their agent is actually ready to ship



SHIP · ±4 CONFIDENCE

02

Make a call on simulation and guardrails capabilities

- Simulation testing is a fast-growing capability and a real gap for AgentScore today
- Only agents already in production, on traces real users emit, can be graded
- Specific situations are hard to test - there's no way to confirm all desired behavior has been exercised
- Tackling this makes AgentScore a well-rounded, complete AI agent testing platform

✓ **Grade what happened**

Real production traces, after the fact - credible, but reactive.

✗ **Simulate before it happens**

Synthetic users, pre-prod - a real gap and a decision point, not yet built.

03

MCP capabilities

- Don't build MCP server conformance testing - already served (MCP Inspector, CLIs), no real moat there
- MCP tool calls are just tool-call spans on an OTel trace
- Ensure ingestion recognizes and labels MCP traces so a report can show "how well did the agent use MCP tools" as a first-class view
- Low-effort, fits the zero-setup story - but not a differentiator alone
- The real moat is pairing this with the AIW hook: AIW agents calling MCP tools get graded automatically

✗ MCP conformance testing: skip

✓ MCP trace labeling: low-effort, do it

✓ MCP + AIW hook: the real moat

04

Fix enterprise readiness

- RBAC, SSO/SAML, and non-VPN access are actual blockers to any external design partner converting
- Demos and closed betas work fine today in the current state
- But these three are table stakes for any long-term engagement



05

Watch item: int4 TrustGate

LOW URGENCY - TRACK, DON'T CHASE

int4 TrustGate is not a head-to-head evaluator, but it's the first to plant an **"AI trust for SAP" flag** - runtime guardrails on agents writing to S/4HANA - inside enterprise accounts Tricentis already sells Tosca into.

It validates the **business outcome** (the SAP document itself) - a framing our behavior-grading doesn't cover - landing right where Tosca's SAP testing strength lives.

TWO IMPLICATIONS

- A A plausible integration/partner surface - grade the agent's behavior, gate its write - more than a rival.
- B If enterprise buyers start expecting "AI trust for my ERP agents," decide whether AgentScore's verdict layer plugs into that story via Tosca.

Brand-new, bootstrapped, unproven - but the most Tricentis-adjacent move in the field. Flag for Tosca / SAP GTM.



We are the **verdict layer** competitors aren't building.

We win by being the grading layer no one else is, **delivered where no one else can reach.**